

Human Pose Estimation

Davide Mattioli, Pablo Jahnen, Kaisar Dauletbek

Introduction

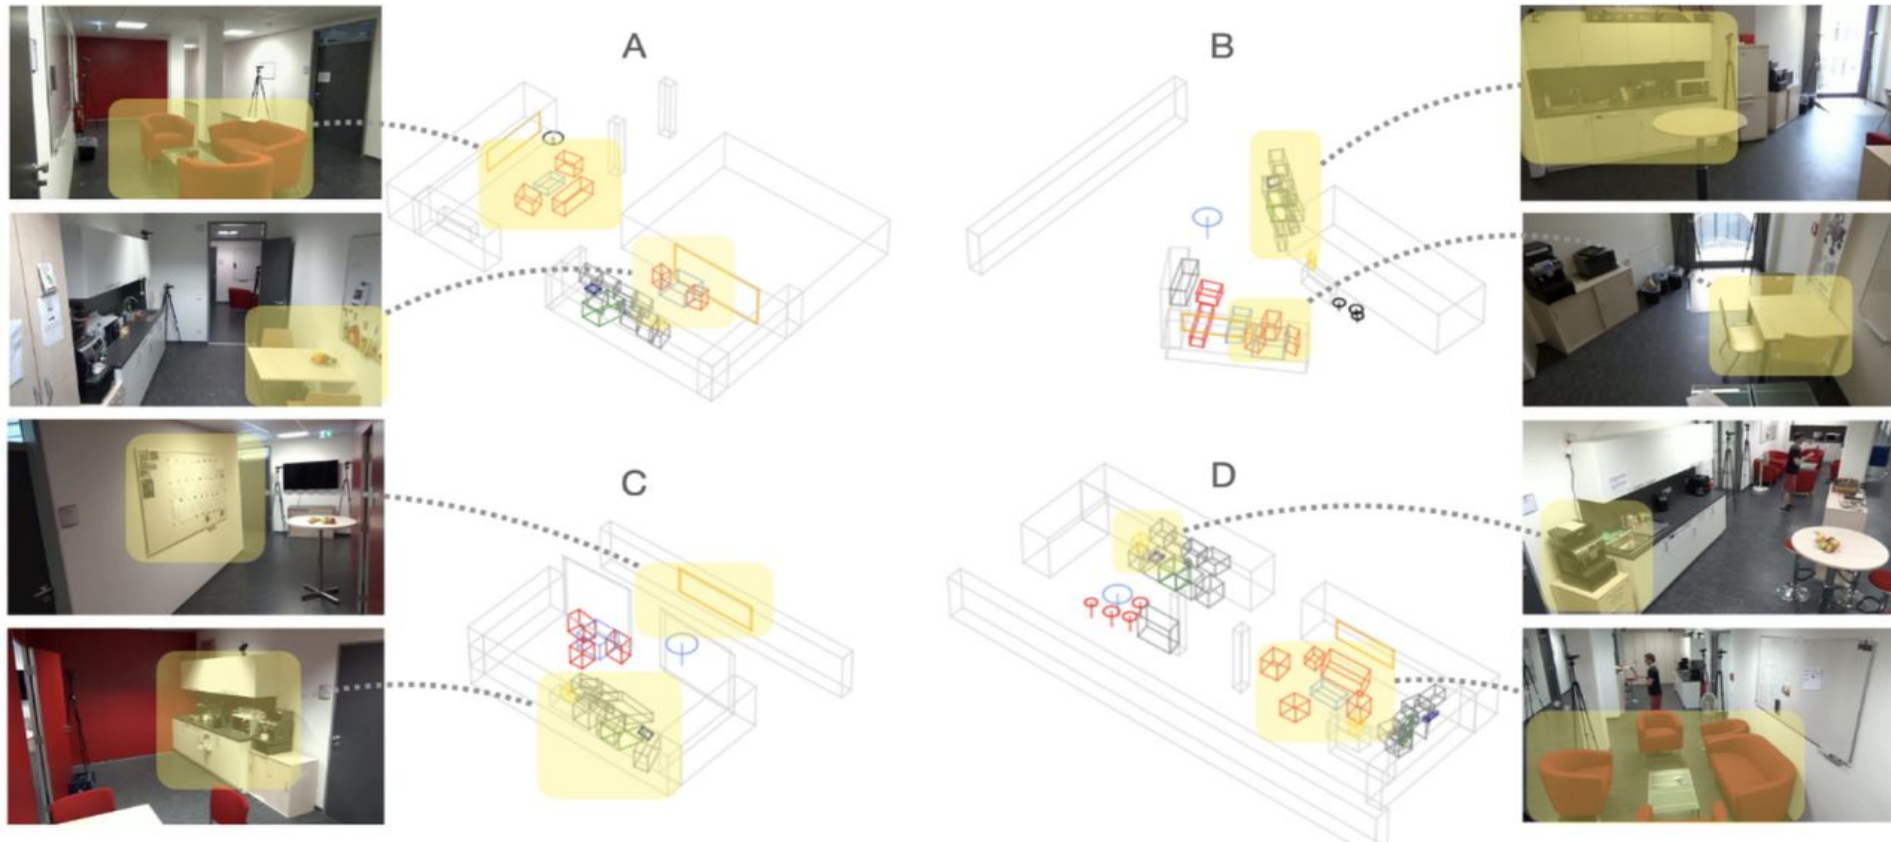
- Objective: Prediction of Poses through graph classification
- Dataset: Humans in kitchens - Realistic behaviour of multiple interacting persons. Data collected in four real office kitchens for a duration between 1.5h and 1.9h
- Graph building: Using 29 keypoint locations that map to the body (nodes), and connecting spatial and temporal neighbors (edges)
- Target: Based on a X frame sequence, we want to predict the activity label of the center frame




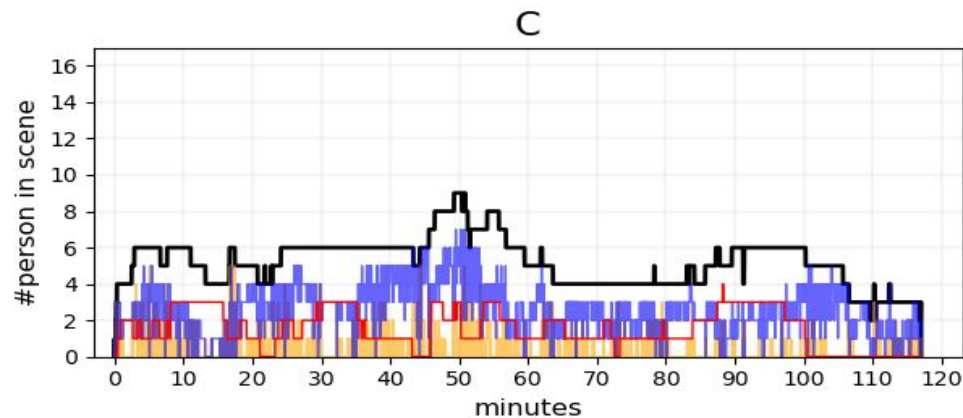
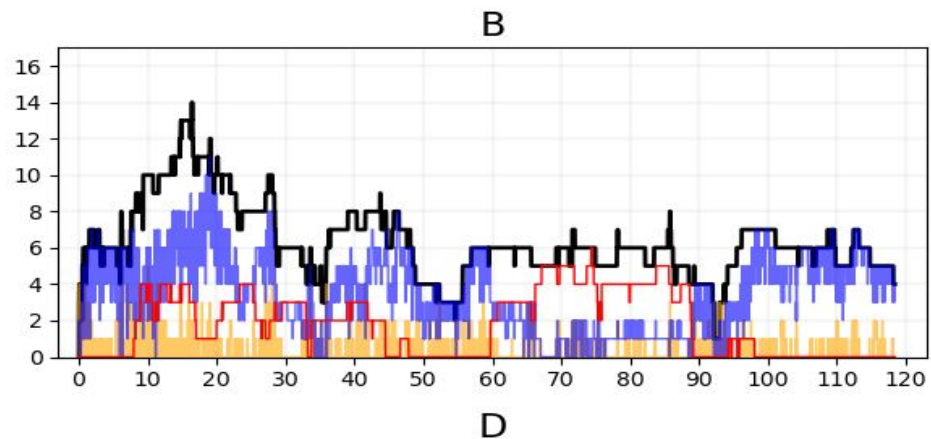
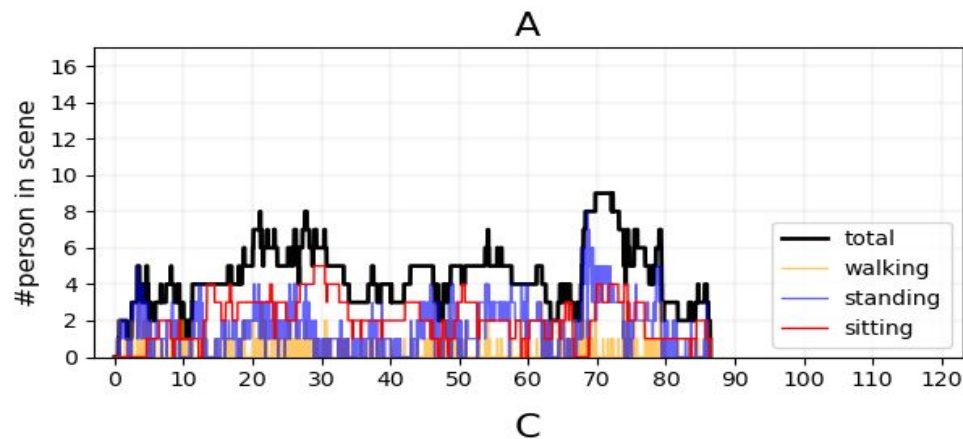
The Dataset

- 4M Annotated Poses of 90 Individuals
- 4 Dataset of different parts of the kitchen

Dataset Regions



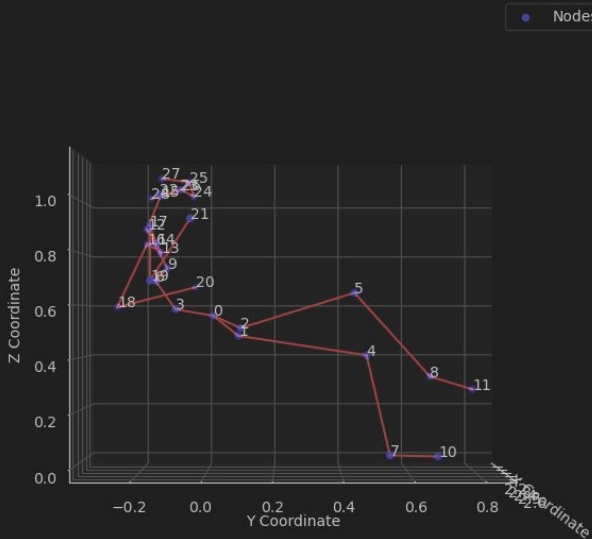
Different Region leads to Different Distributions



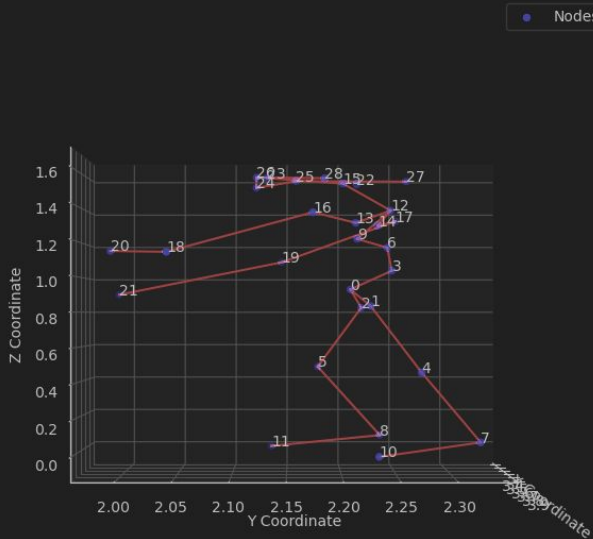
Poses Examples

Each node has (X,Y,Z) coordinates

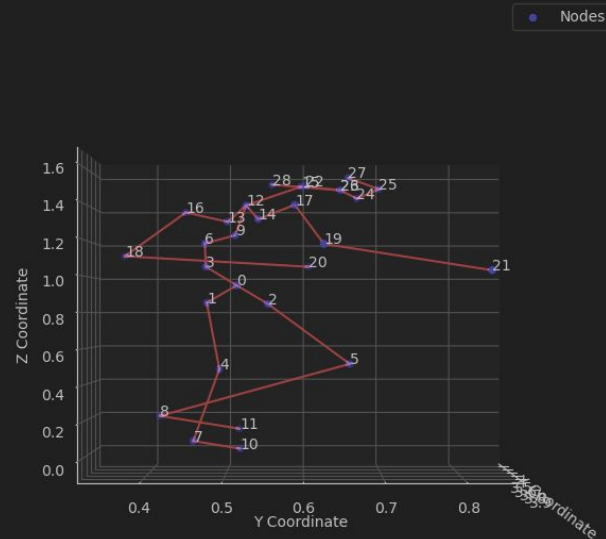
Person Plot ID =1



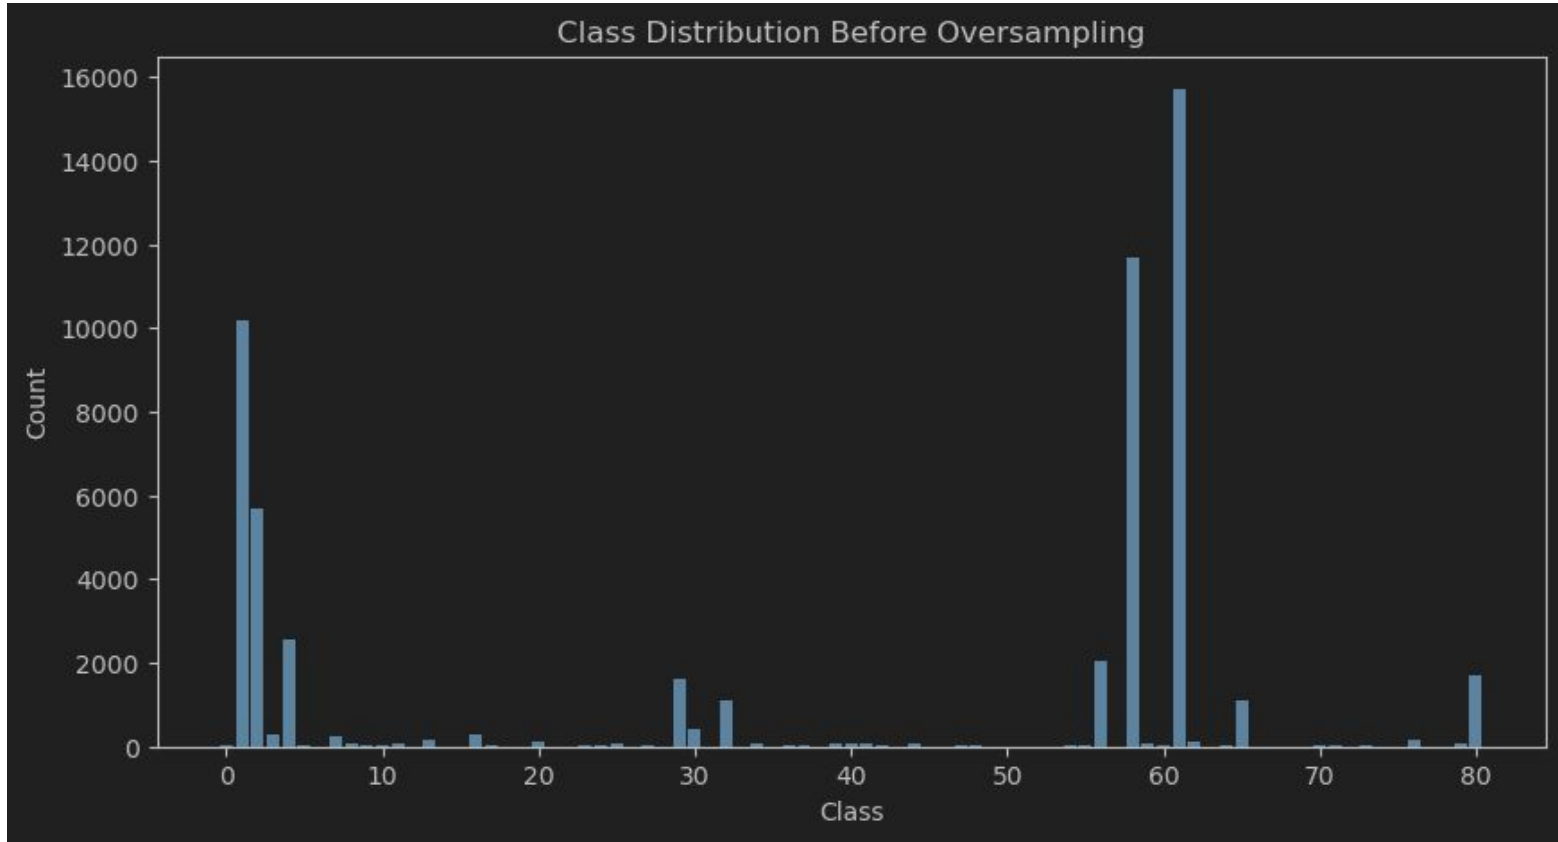
Person Plot ID =2



Person Plot ID =3



Class Distribution



Architectures

GCM

GINE


GINE with edge features

Graph Transformer

GCNConv Layer

Message passing: Combines node features and normalizes by degree

Node update: sparse matrix multiplication to aggregate messages without additional scaling mechanisms.



Gine Layer

Message passing: the GINELayer combines source node features and transformed edge features in message computation.

Node update: Uses a learnable trainable epsilon to scale the contribution of node features during updates.

```
src, dst = edge_index

# updating edge features with mlp, bringing them to the same dimension as node features
edge_attr = self.edge_mlp(edge_attr.float())
message = x[src] + edge_attr

# aggregating messages
aggr_out = scatter_mean(message, dst, dim=0, dim_size=x.size(0))

# node update
out = self.node_mlp((1 + self.eps) * x + aggr_out)
```

2-Step- Gine

Uses an additional pre-trained sit-stand model to compute extra input features for the GINE model.

```
sit_stand_model = GINE(  
    num_features=3,  
    num_classes=3,  
    hidden_channels=256,  
    num_layers=8,  
    use_batch_norm=True  
) .to(device)  
  
model = GINE(  
    num_features=config["input_dim"],  
    num_classes=output_dim,  
    hidden_channels=128,  
    num_layers=config["num_layers"],  
    use_batch_norm=True  
) .to(device)
```

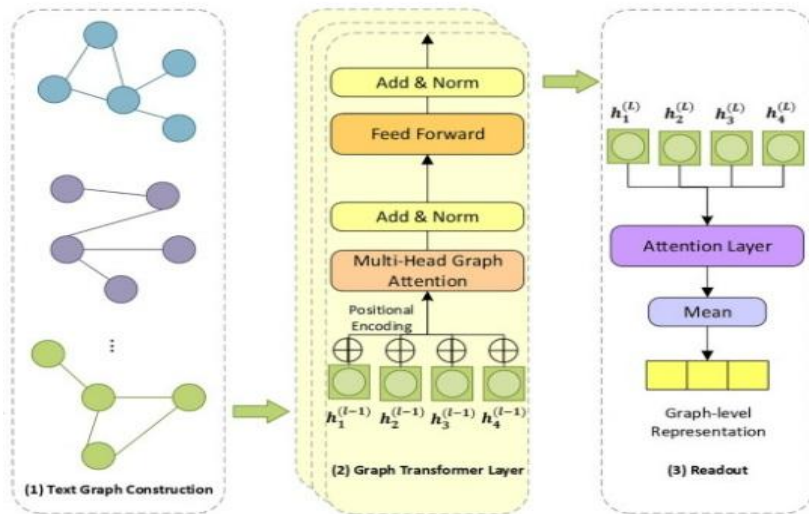
The pre-trained model predicts the sit-stand state, which is added as an additional feature column to the input node features

```
sit_stand_pred = torch.argmax(sit_stand_out, dim=1)  
sit_stand_pred = torch.repeat_interleave(sit_stand_pred, 145, dim=0).unsqueeze(0).T  
batch.x = torch.cat((batch.x, sit_stand_pred), 1)
```

Graph Transformer

Message passing: Each TransformerConv layer performs message passing based on self-attention

Self-attention: dot-product attention mechanism but applied locally in the graph neighborhood




Experimental results – GCN Baseline

256 hidden size, 2 layers.

Test Scores:

- F1: **0.0513**
- Multiclass accuracy: **0.6542**




Experimental Results - Just a Big Transformer

200 epochs, 256 hidden size, 16 heads.

Tweaks: batch normalization across features

Test F1 Scores:

- Without learning rate scheduler: **0.3775**.
- With cosine LR scheduler: **0.3737**.
- Adding feature normalization across batch: (Result pending completion).




Experimental Results - GINE

Graph Isomorphism Network with Edge Features (GINE)

Tweaks: +plateau scheduler, + edge features, + undersampling

- **Configuration 1:**
 - 64 hidden size, 4 layers
 - **Test F1 Score: 0.2983.**
- **Configuration 2:**
 - 512 hidden size, 12 layers, same epochs and batch size.
 - Similar performance to Configuration 1.



Experimental Results - 2-Step GINE

Two-Step Classification

First, we trained a small GINE to classify sitting, standing, and none of that (trained on 20% of the data):

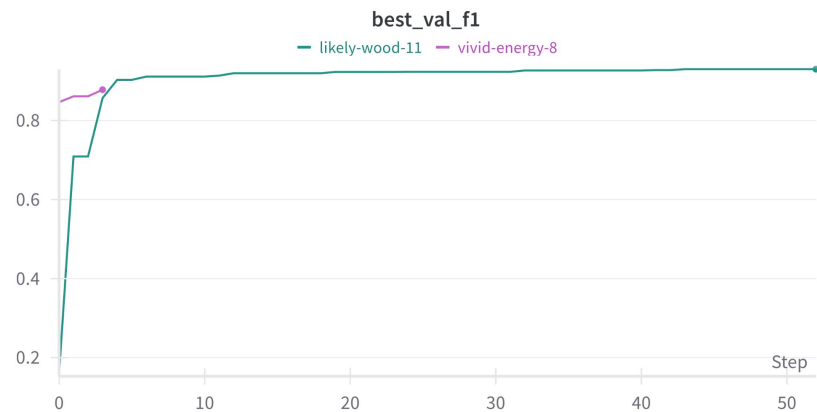
- Additional feature dimension for `sit_stand` label.
- Configuration: 256 hidden size, 4 layers, batch size, early stopping.
- **Test Scores:**
 - F1 Score: **0.9302**

Then, we trained a larger model for which the features were augmented by the prediction of the sit/stand model.


- Additional feature dimension for `sit_stand` label. (Did it in 2 ways)
- Configuration: 512 hidden size, 8 layers, 1028 batch size, early stopping.
- **Test Scores:**
 - F1 Score: **0.2489**.
 - Multiclass Accuracy: **0.6012**.

Notes: GINE is small, quick to train

Sit/stand:



Main Model:



Experimental Results - Back to Transformer

Lastly, we trained a smaller transformer model, but added the following tweaks:

Tweaks: + class weights to the loss function

First, we trained a small GINE to classify sitting, standing, and none of that:

- Additional feature dimension for `sit_stand` label.
- Configuration: 256 hidden size, 2 layers, 4 attention heads.
- **Test Scores:**
 - F1 Score: **0.0934**
 - Multiclass Accuracy: **0.8032**

Discussion

What we could have done also:

Experiment more with spatio-temporal features

Positional encodings (RW in space, RW in time)

Add weights to the metrics